


The Battle Card Engine

A daily operating system that turns scattered relationship signal into focused human action.



3–5x

MORE ACCOUNTS QUALIFIED & NURTURED / WEEK

+35%

BETTER PROSPECT CONNECT RATE

~8–10%

EST. CHURN REDUCTION, LAST 2.5 MONTHS

Nikhil More

Performance Marketing Operator → AI Generalist

NSDC AI Generalist Certified

19 Years B2B Performance Marketing

\$55M+ Pipeline Generated

THE BOTTLENECK


Prioritisation, not personalisation, is the hard part.

The point is not to automate outreach. It is to concentrate human attention where it can matter.

Every "AI in sales" post is about writing a better cold email. That is not the bottleneck anymore — everyone has AI drafting outreach now. The real bottleneck is earlier: knowing which account is worth a human's time today, before anyone opens a tab to check.

I built a system around that decision. It maps the account, reads relationship signal, scores what changed, and produces a practical battle card. The output is not more data. It is a next action.

THE ATTENTION ALLOCATION FUNNEL

**3–5x**

MORE ACCOUNTS QUALIFIED & NURTURED

Self-tracked performance range

+35%

BETTER PROSPECT CONNECT RATE

Compared with earlier approach

~8–10%

ESTIMATED CHURN REDUCTION

Over the last 2.5 months

THE SIGNAL-TO-SEQUENCE FRAMEWORK

Four layers. One useful decision.

Adapt this to whatever stack you have. The architecture matters more than the specific tool.

01 Discover

Map the account before you target a person. Find decision-makers, influencers and users — not just the one contact in the CRM.

02 Listen

Pull signal from every channel the relationship actually lives in: research, email, calls, WhatsApp and meeting notes.

03 Prioritise

Do not act on every signal. Score changes and surface only the accounts that earned attention today.

04 Brief

Hand over a decision, not a dashboard: who to contact, why now and what to open with.

WHY PRIORITISATION IS THE HARD PART

Most teams already have more data than they can use. The scarce resource was never information; it is knowing which three accounts, out of two hundred, actually deserve five minutes today.

LAYER 1

Map the buying committee before you target it.


Account research becomes a standing input to the campaign, rather than an hour of manual work.

THE PROBLEM THIS REPLACED

Before this, targeting a new account meant manually building the persona map by hand — searching LinkedIn, cross-checking titles, and guessing who actually influenced a buying decision versus who merely appeared in the org chart. It was slow and half-informed even when it was done.

WHAT IT DOES

Given a campaign, the system scrapes B2B intelligence through Firecrawl running against ZoomInfo- and Apollo.io-style sources. It finds the buying committee — decision-makers, influencers and end users — separately labelled, so outreach does not guess which contact matters.



Illustrative account names and roles only — no client data shown.


3–5x more accounts qualified and nurtured per week

Self-tracked range. The gain is volume: one operator can now cover mapping work that earlier required a research pass per account.

HOW IT WORKS

From campaign brief to a scored account.

The system removes the first hour of manual account research, but leaves relationship judgement with the operator.




HOW IT IS BUILT

A campaign brief goes in. The scraper pulls account and contact-level data, an AI layer classifies each person into a role, and the output is a scored account with a persona map attached — ready for the next layer to act on.

WHAT CHANGED

Account targeting moved from a manual research task to a standing input. I am not spending the first hour of a campaign figuring out who to talk to; that work is already done when I open the account.

ILLUSTRATIVE READINESS SCORE



The account map turns a campaign from a research exercise into a ready-to-work decision set.

LAYER 2

Keep the relationship in one record.

A signal should not evaporate simply because it happened outside the official channel.


THE PROBLEM THIS REPLACED

Account intelligence used to live in five disconnected places: my inbox, my phone, a WhatsApp thread, handwritten notes from an in-person meeting, and my memory of a call from three weeks ago. None of it made it into one place unless I manually typed it in, which mostly did not happen.

WHAT IT DOES

The system reads and logs signal across every channel the relationship actually happens in: buying signals, intent data, email content, phone-call intelligence and WhatsApp or messenger activity. For face-to-face meetings, a voice path turns a spoken recap into a transcript in the same record.

SEVEN-DAY SIGNAL TIMELINE — FICTIONAL ACCOUNT



ACCOUNT RECORD

One connected picture of the relationship

+35% better prospect connect rate

Because the first message can be based on what actually happened across the relationship. Tracked comparison; not an audited study.

HOW IT WORKS

Capture the context. Keep the human judgement.

The system is a memory layer for my own relationships — not a substitute for judgement or a surveillance tool.

IMPORTANT BOUNDARY


This is internal intelligence capture on my own sales conversations and relationships, feeding my own follow-up decisions — not third-party surveillance. It exists so that a signal from a WhatsApp message is not lost because it did not happen in the official channel.

HOW IT IS BUILT

Each channel feeds into the marketing automation system as a logged event tied to the account record. The voice-agent path adds a transcription step before the same thing happens. Nothing here replaces judgement; it stops relevant context from evaporating.

WHAT CHANGED

The account record stopped being whatever I remembered to write down and became a more complete picture, regardless of which channel the conversation used.



The account record becomes a complete picture.

The recommended message is informed by the relationship, not a template that ignores it.

LAYER 3

Do not surface every signal. Surface the change that matters.

More data does not make prioritisation easier. It makes the pile larger unless the system applies triage.

THE PROBLEM THIS REPLACED

Every morning started with the same question: which of these accounts do I actually work today? With more accounts and more channels of signal, that question got harder, not easier. More data simply created a larger pile to sift through by hand.

WHAT IT DOES

The system scores account activity and only surfaces accounts where something changed today: a buying signal fired, a contact engaged, a competitor mention appeared, or a stage-relevant event happened. It tells me which handful earned attention before I open anything.

ILLUSTRATIVE PRIORITY SCORE — NORTHSTAR SYSTEMS

Intent 40 pts	Engage 25 pts	Recency 20 pts	Ctx 15 pts
------------------	------------------	-------------------	---------------

RELEVANCE THRESHOLD

92 / 100 — Priority 01

Human can accept, ignore or override the recommendation.

Prioritisation, not personalisation, is what moved the numbers.

Point the work at the right five accounts instead of spreading it evenly across two hundred.

THE DAILY WORKFLOW

From noisy pipeline to five battle cards.

This is the differentiation layer: the system makes an attention recommendation, then gives the operator the context to judge it.

HOW IT IS BUILT

Every signal logged in Layer 2, plus the account data mapped in Layer 1, feeds a scoring pass. Accounts that cross a relevance threshold are surfaced; the rest wait. It turns "we have a lot of data" into "here is what matters now".

WHAT CHANGED

I stopped starting my day by reviewing a pipeline dashboard and started by reading five battle cards. The decision about where to spend the morning is made before I sit down.

AN HONEST LIMIT

The score should not pretend to know everything. It is a recommendation based on visible signals and needs an operator who understands the account and can challenge the output.

TODAY — ACCOUNTS THAT MOVED

- 01 Northstar Systems**
Pricing engagement after competitor mention
- 02 Crestline Cloud**
Champion replied to nurture
- 03 AxisWorks**
Stage-relevant event added
- 04 Brightline**
No significant change

The operator sees a small, defensible work list — not a dashboard that still requires manual sorting.

LAYER 4

Close the gap between "what happened" and "what now".

A useful output should reduce preparation time and make the next action obvious.

THE PROBLEM THIS REPLACED

Even after knowing which account to work, there was still a research step before every call or email: re-reading old notes, checking what was last discussed, and finding an opening line. That was five to ten minutes of prep repeated for every touch.

WHAT IT DOES

For every account the prioritisation layer surfaces, the system pre-builds a battle card: what happened, why this account is worth today's attention, and a suggested opening for outreach. It also tracks pipeline stage and produces a weekly report with the recommended next step.

BATTLE CARD — FICTIONAL ACCOUNT

Northstar Systems

PRIORITY 01

HIGH INTENT

WHY NOW

Pricing engagement after a competitor mention.

OPEN WITH

Ask whether campaign expansion changed the priority.

NEXT ACTION

Call VP Marketing today. Log the outcome.

The output is never raw data.

It is a next action: call this account, here is why, and here is how to open.

HOW THE BRIEF CREATES ACTION

A battle card for today. A pipeline view for the week.

The same logic works at two levels: a ready-to-act account brief and a portfolio view of what should move next.

HOW IT IS BUILT

The battle card pulls directly from the account record built in Layers 1 to 3 and formats it as a decision, not a data dump: call this account, here is why, and here is what to open with. The weekly pipeline report uses the same logic at portfolio level — stage, reasoning and next action.

WHAT CHANGED

The gap between "the system told me something happened" and "I know what to do about it" closed. Across every layer, the design principle remains the same: the output is not raw data; it is a next action.

WEEKLY PIPELINE — ILLUSTRATIVE

- Discovery**
Confirm buying committee
- Evaluation**
Share proof point
- Proposal**
Unblock decision process

Prep per account touch dropped from a research task to a read.

Weekly pipeline review became a five-minute read instead of an hour of manual report pulling.

THE REST OF THE STACK

The ABM engine is not the only system running.

The common design principle is simple: find the repeat cost and build the thinnest system that removes it.

Inbox and Calendar Assistant


It triages my inbox — separating what is informative from what needs a reply — drafts responses from historical context, and leaves them in drafts for review. For recurring meetings, it builds a brief from prior notes and generates the artifact I need: a deck, research, ad copy or a project plan.

Meeting-to-Action Flow

It captures notes and attendees, produces minutes, then pushes tasks and deadlines into Asana from the transcript, assigned to the right people. It can also turn Asana activity into productivity and performance reports tied to KRIs.

Research and Planning Agents

A research agent tracks industry trends and surfaces topics relevant to my work and audience. A lower-stakes holiday planner compares destinations, builds a budget and creates a practical itinerary.



Unstructured human activity becomes structured downstream action.

Across every system, the human remains responsible for the decision.

WHAT IT BOUGHT BACK

Time to think about strategy, not only execute it.

The goal was never to run AI for its own sake. The goal was to buy back time from repeat coordination.

None of this was built for its own sake. Between this system and the others, I am getting back an estimated 20+ hours a week. That is a self-tracked number, not an audited one, but it has held steady for long enough that I trust it.

What matters more than the number is what that time is now used for: more time to think about strategy instead of executing it, and more time to build the next system instead of doing the manual version of the current one.



WHERE THE CAPACITY GOES

Strategy and planning



Improving active systems



Building the next system



THE OPERATING BELIEF

Do not begin with: what can AI do? Begin with: what manual decision have I made 40 times already?

A PRACTICAL THESIS

I stopped asking AI to write better emails. I asked it to tell me who to email.

That question — what should I actually work on right now? — is the one worth building for. Everything in this field note is one answer to it. It will not be the last system I build to answer it.

**IF YOU ARE BUILDING SOMETHING SIMILAR**

Happy to compare notes on the repeat cost, the workflow and the thinnest version worth shipping first.